

Video Transcript - Part 7

Slide 1

Hey everyone, welcome to part seven. This time we're looking at Apache Spark, and specifically how it runs on Amazon EMR. Spark is one of the most widely used big data processing engines out there, and it powers a lot of machine learning pipelines too, and EMR is how AWS lets you run it without managing the underlying cluster yourself.

Slide 2

Here's the plan. We'll start with what Apache Spark actually is, then look at why it's so much faster than older approaches, walk through its basic architecture, move on to why you'd run Spark on Amazon EMR specifically, cover the key features and benefits, and finish with some common use cases.

Slide 3

So what is Apache Spark? It's an open source, distributed processing engine built for big data. It spreads both data and computation across a cluster of machines, so large jobs get split into smaller pieces that run in parallel. What makes Spark especially popular is that it's general purpose, the same engine handles batch processing, real time streaming, SQL queries, machine learning through its MLlib library, and even graph processing. And it supports multiple languages, so you can write Spark code in Python, Scala, Java, R, or plain SQL.

Slide 4

One of Spark's biggest selling points is speed. Traditional MapReduce, which is Hadoop's original processing model, writes intermediate results to disk between steps, which is reliable but slow. Spark, by contrast, keeps data in memory across those steps whenever possible, only touching disk when it truly has to. For iterative workloads, like training a machine learning model over many passes through the data, that difference in speed can be dramatic, often ten to a hundred times faster.

Slide 5

At a high level, Spark's architecture has three parts. The Driver Program coordinates the whole job, it splits your code into tasks and decides how work gets distributed. The Cluster Manager handles allocating resources, deciding which machines get which tasks. And the Executors are the actual worker processes running on those machines, they execute the tasks and hold data in memory for fast access. Together, these three pieces are what let Spark scale a single job across dozens or even thousands of machines.

Slide 6

Now, setting up and managing your own Spark cluster means provisioning servers, installing and configuring Spark yourself, and constantly babysitting capacity. That's exactly where Amazon EMR comes in. EMR is AWS's fully managed big data platform, and it takes care of provisioning, configuring, and tuning Spark clusters for you. It reads and writes data directly from Amazon S3, which decouples your storage from your compute, meaning your data sticks around even after you shut a cluster down, and you only pay for compute while a job is actually running. And EMR gives you flexible deployment options too, you can run Spark on EC2 for full control, use EMR Serverless if you don't want to manage clusters at all, or run it on EMR on EKS if you're already using Kubernetes.

Slide 7

A few features make Spark on EMR genuinely convenient. It auto-scales, growing or shrinking your cluster automatically as workload demands change. It supports spot instances, which can cut costs significantly for jobs that can tolerate some interruption. EMR Studio and EMR Notebooks give you an interactive, Jupyter based environment for developing and testing Spark code directly against your cluster. And it integrates deeply with the rest of AWS, things like the Glue Data Catalog for metadata, Athena for querying, SageMaker for machine learning, and CloudWatch for monitoring.

Slide 8

In practice, you'll see Spark on EMR used for a handful of common scenarios. Large scale ETL, cleaning and transforming massive datasets before they land in a warehouse or data lake. Streaming analytics, processing real time data as it arrives using Spark Streaming. Machine learning pipelines, doing feature engineering and training models at scale with MLlib. And interactive data exploration, where data scientists run ad hoc analysis directly through notebooks.

Slide 9

So to recap, Apache Spark is a fast, general purpose, distributed processing engine that keeps data in memory to outperform older disk based approaches. And Amazon EMR is how you run Spark, and the rest of the big data ecosystem, as a fully managed service on AWS, without provisioning or babysitting the cluster yourself. Thanks for watching, and I'll see you in the next one.